



AI for Augmented Insurers.

Detect, explain & correct data drift in a machine learning system.

November 2021

Agenda

1. Introduction
2. Detect data drift
3. Explain data drift
4. Correct data drift
5. Demo
6. Conclusion



Introduction



Context

- We consider a ML system answering a binary classification problem where:
 - $Y \in \{0, 1\}$: target
 - $X \in \mathcal{X} \subset \mathbb{R}^d$: input
- Let $\hat{f}(x)$ a model which answers the binary classification problem and is already trained on some dataset.

Some idea developed in the presentation requires that \hat{f} is a tree-based model (Gradient Boosting (XGBoost, LightGBM), Random Forest, CART, etc.).

- Let $\mathcal{D}_1: (X_{i,1}, Y_{i,1})_{i=1, \dots, n_1}$ (i.i.d.) and $\mathcal{D}_2: (X_{i,2}, Y_{i,2})_{i=1, \dots, n_2}$ (i.i.d.) two datasets corresponding to the binary classification problem. Let P_{X_1, Y_1} and P_{X_2, Y_2} the corresponding distributions.

Our goal is to **study the data drift between \mathcal{D}_1 and \mathcal{D}_2**

- In a typical situation \mathcal{D}_1 is the training/validation data and \mathcal{D}_2 is the production data.

In practice depending on the use case, $(Y_{i,2})_{i=1, \dots, n_2}$ may not be observed.

What is data drift ?

Definition: A data drift between \mathcal{D}_1 and \mathcal{D}_2 corresponds to the case where $P_{X_1, Y_1} \neq P_{X_2, Y_2}$

Two categories of data drift:

- Case 1: **Covariate shift**
 - the conditional distributions are the same: $P_{Y_1|X_1} = P_{Y_2|X_2}$, but the distribution of inputs changes: $P_{X_1} \neq P_{X_2}$
- Case 2: **Concept drift**
 - the conditional distributions are not the same: $P_{Y_1|X_1} \neq P_{Y_2|X_2}$

Property: P_X and $P_{Y|X}$ characterize the distribution $P_{X,Y}$

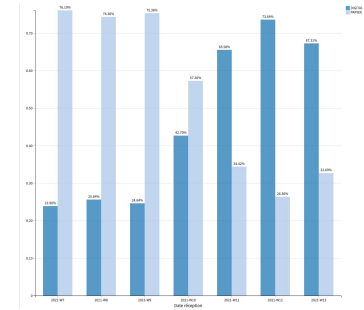
proof: $dP_{X,Y}(x, y) = dP_{Y|X}(y|x) \cdot dP_X(x)$

Property: These 2 cases cover all possible cases

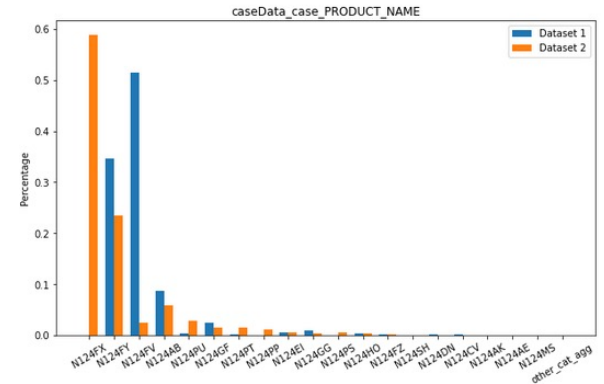
proof: In covariate shift, if P_X is unchanged, then $P_{X,Y}$ is given by P_X and $P_{Y|X}$ which are both unchanged. Hence $P_{X,Y}$ is unchanged -> no data drift

Case 1: Examples of covariate shift

- The distribution of some input variable that changes with time. (e.g. evolution of inbound channel repartition: more digital, less paper)
- Some input variable that may take a new value (e.g. new product, new option, new functionality, etc.)
- Hand-based input with a user error in it (e.g. you're expecting an input to be in M\$, and it's in \$ instead)



Evolution of inbound channel repartition



Case 2: Examples of concept drift

- **Censoring of some observations**

- At training time, **we filter the data to keep only the samples for which we know the target Y** . Let $D \in \{0, 1\}$ be the censoring variable. Training data is (X_i, Y_i, D_i) with $Y_i = \infty$ if $D_i = 0$ (training data has the distribution $P_{X,Y|D=1}$).
→ Typical situation is the **presence of non stated cases** (accepted, rejected and unknown status).

- **Non-stationary environment**

- Change of the regulation (e.g. credit agreement).
- Important events that change behaviors (e.g. propensity to travel during Covid crisis).

- **Spatial drift**

- Model trained on French insurance data and deployed in the Italian market.



Detect data drift

Measure of drift - Numerical variables

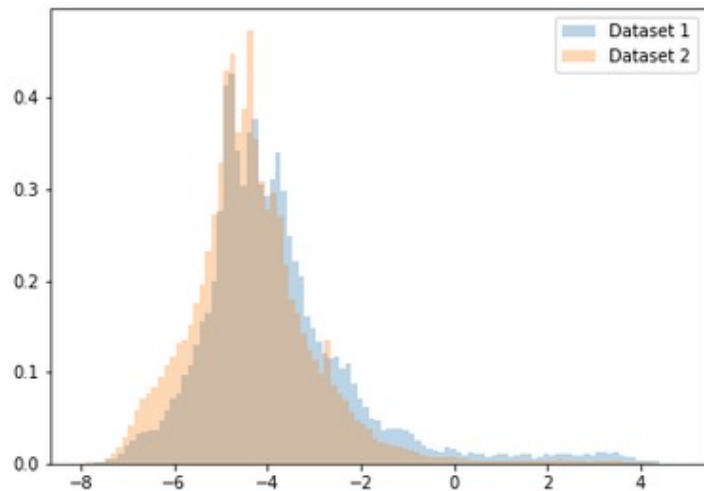
- Difference of means.
- Wasserstein distance

$$W(V_1, V_2) = \int_{-\infty}^{+\infty} |F_{V_1}(v) - F_{V_2}(v)| dv,$$

with $F_V(t) = P(V \leq t)$ the CDF of V .

- Kolmogorov-Smirnov 2 sample test.

Distribution of variable V



```
{'mean_difference': -0.514477079781603,  
  'wasserstein': 0.5144829964200954,  
  'kolmogorov_smirnov': KstestResult(statistic=0.13030577961608247, pvalue=0.0)}
```

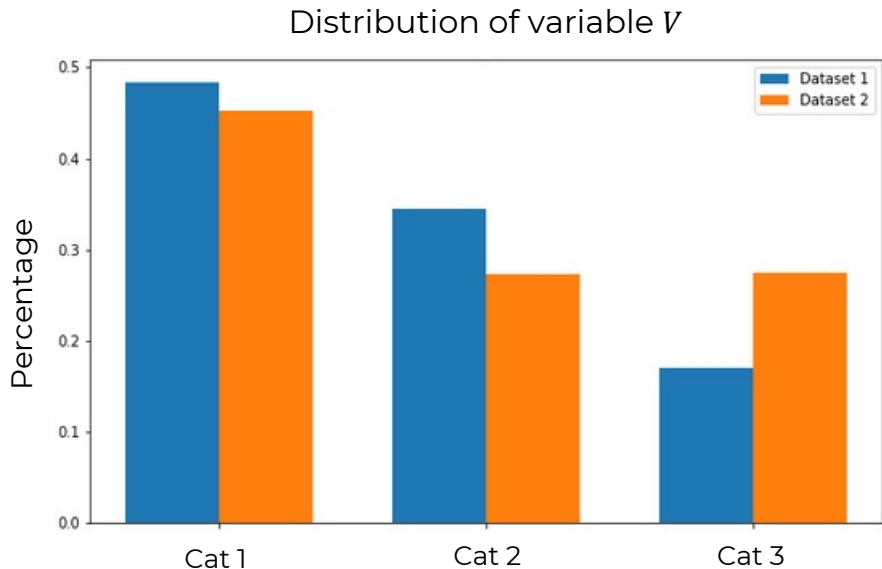
Measure of drift - Categorical variables

- Wasserstein distance (we assume distance between 2 categories is equal to 1*).
- Chi2 test

Contingency table :

	Cat 1	Cat 2	Cat 3
X1	1152.0	821.0	407.0
X2	1909.0	1154.0	1163.0

* This corresponds to Wasserstein distance between dummy representations, with $\| \cdot \|_{\infty}$ norm

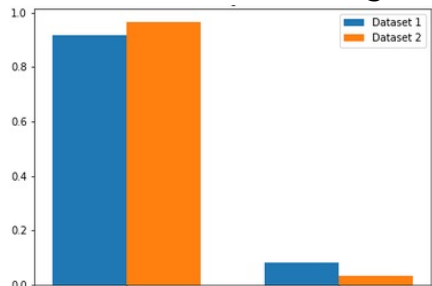


```
{ 'wasserstein': 0.10419273246449548,  
  'chi2_test': { 'chi2_stat': 99.2937, 'p_value': 2.74556e-22} }
```

What indicators do we track in the ML system ?

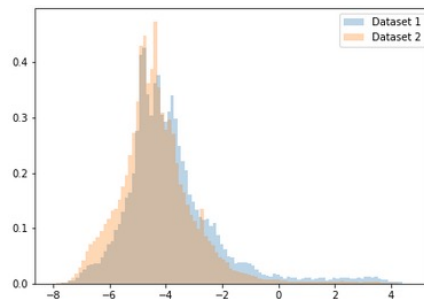
- Distribution of predictions of the model
- Distribution of the target
- Performance metrics

Distribution of the target



```
{'wasserstein': 0.04893272477192796,  
'chi2_test': {'chi2_stat': 1655.184463874216,  
'p_value': 0.0,
```

Distribution of predictions



```
{'mean_difference': -0.514477079781603,  
'wasserstein': 0.5144829964200954,  
'kolmogorov_smirnov': KstestResult(statistic=0.13030577961608247, pvalue=0.0)}
```

Performance metrics

```
log_loss valid: 0.17342663278191264  
log_loss prod: 0.10822475472437297  
AUC valid: 0.8950968405582416  
AUC prod: 0.8393885465363519
```

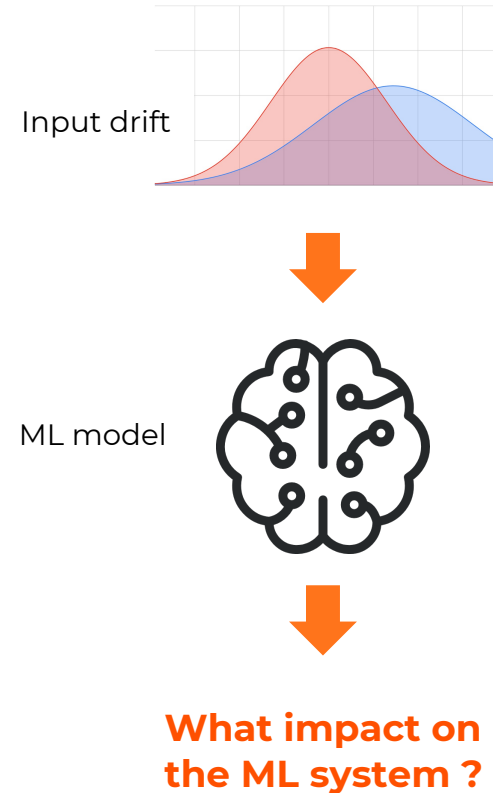
We focus on data drift that have an impact on the ML system (we don't track data drift on all inputs of the model)



Explain data drift

Approach 1: Model based approach

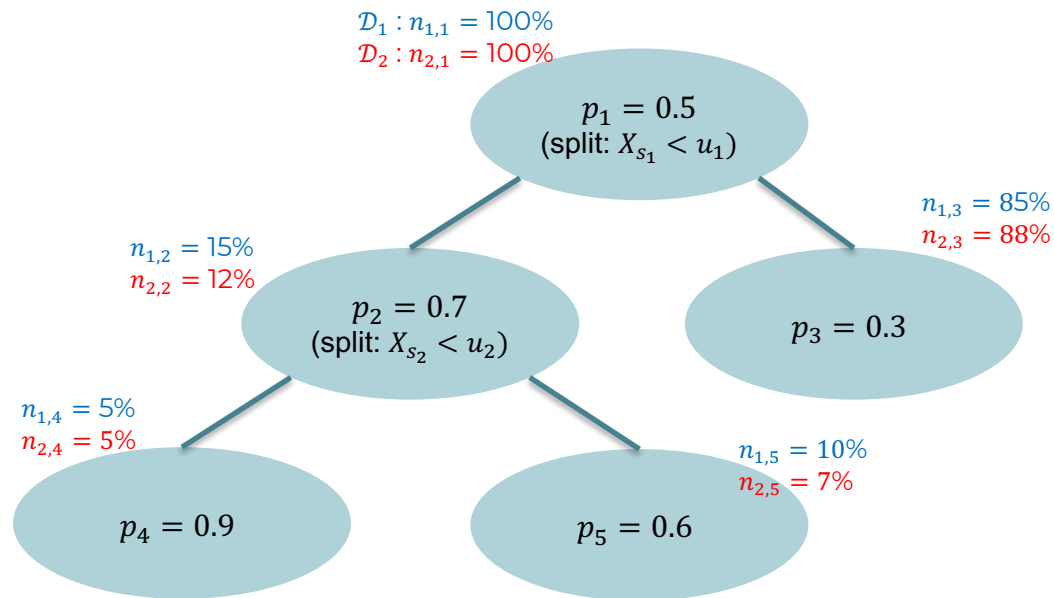
- Ideas:
 - What is the impact of the data drift of a given input on the ML system ?
 - Compute **drift values** of each input (i.e. contribution of the feature to the global data drift)
 - Importance of a data drift on some input of the model should be lowered if the given input is not important in the model
 - Need to study the data drift through the lens of the model
 - **Model specific approach** (here for tree-based model)



Drift values: Calculus ½ (individual tree)

Definition:

- $n_{i,j}$ ($i = 1, 2$; $j = 1, \dots, K$) : proportion of samples of dataset i in node j
- p_j ($j = 1, \dots, K$) : predicted value associated to node j
- s_j, u_j ($j = 1, \dots, K$) : feature index and value used for split j ($s_j, u_j = -1$ if node j is a terminal leaf)
- l_j, r_j ($j = 1, \dots, K$) : indexes of the left and right child nodes of node j ($l_j, r_j = -1$ if node j is a terminal leaf)



Drift values: Calculus 2/2

- Different measure of **split contribution** $S_j (j = 1, \dots, K)$ to the data drift:

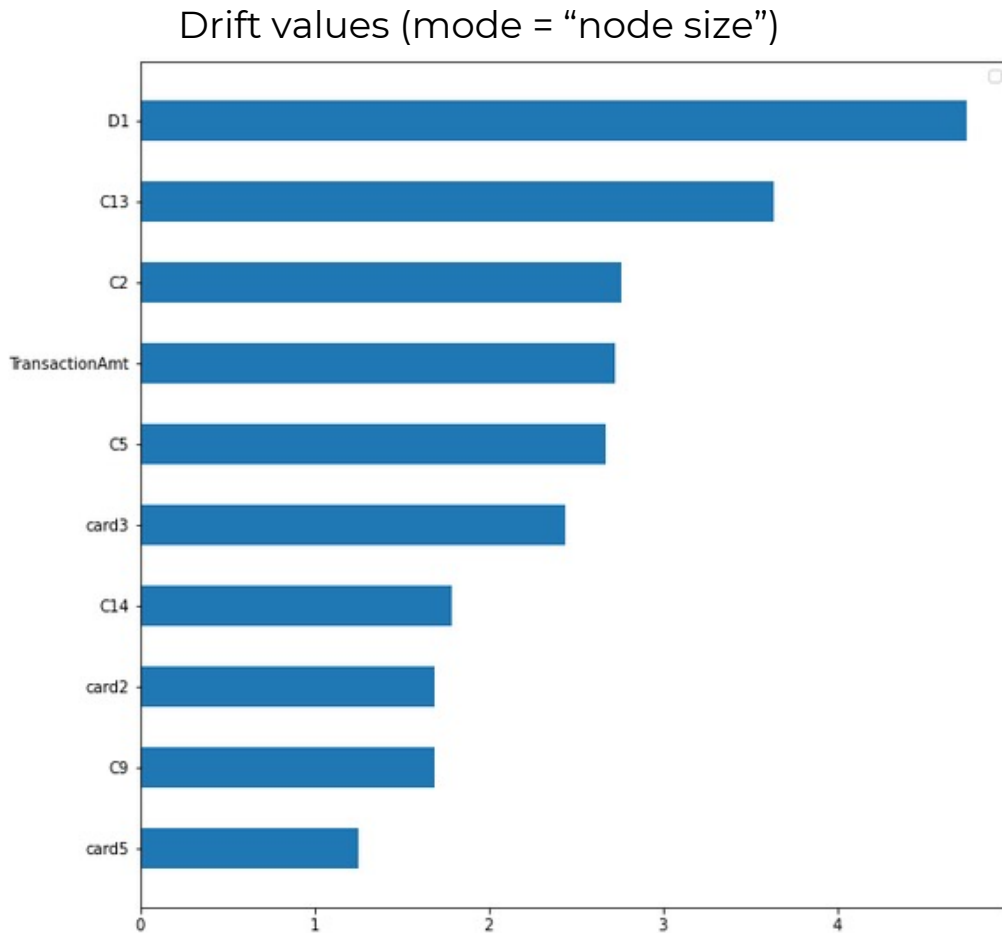
- node size: $S_j = \left| \frac{n_{2,l_j}}{n_{2,j}} - \frac{n_{1,l_j}}{n_{1,j}} \right| * \min(n_{1,j}, n_{2,j})$
- mean: $S_j = c_{2,j} - c_{1,j}$, with $c_{i,j} = (n_{i,l_j} * p_{l_j} + n_{i,r_j} * p_{r_j} - n_{i,j} * p_j)$
- mean with normalization: $S_j = \left(\frac{c_{2,j}}{n_{2,j}} - \frac{c_{1,j}}{n_{1,j}} \right) * \min(n_{1,j}, n_{2,j})$

(in all case we set $S_j = 0$ if $l_j = -1$)

- Based on the split contributions, we can compute **drift values**. Let T be number of trees in the model and s_j^t, K^t, S_j^t be the above values for the tree t . The feature contribution of the feature k ($k = 1, \dots, d$) is defined as:

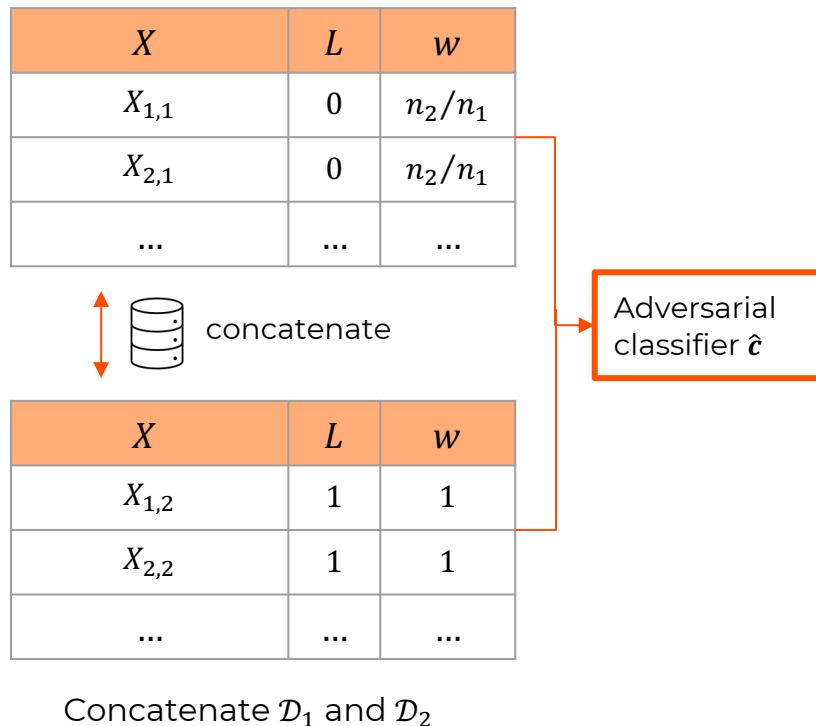
$$F_k = \sum_{t=1}^T F_k^t, \text{ with } F_k^t = \sum_{j=1, \dots, K^t / s_j^t = k} S_j^t$$

Drift values: Example



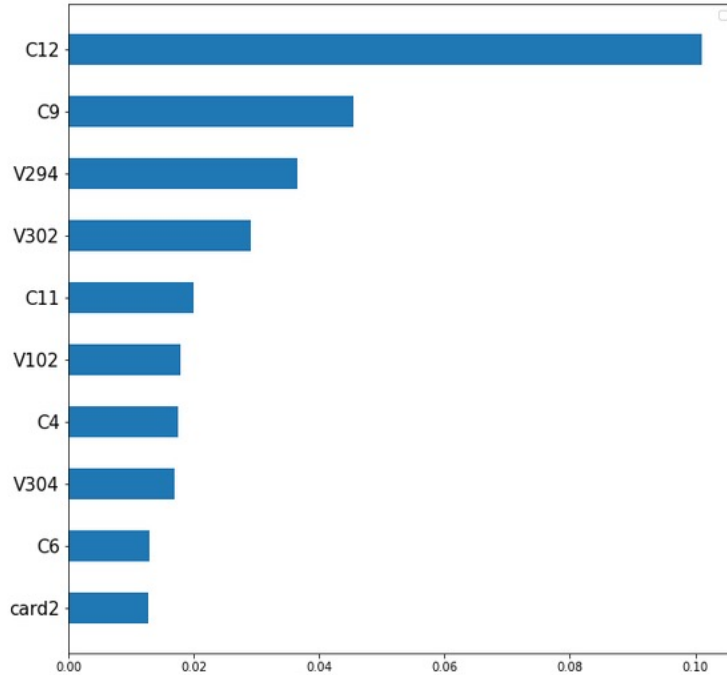
Approach 2: Adversarial approach

- Consider $(X_{i,1})_{i=1,\dots,n_1}$ and $(X_{i,2})_{i=1,\dots,n_2}$ and let:
 - $L_{i,1} = 0, w_{i,1} = n_2/n_1$
 - $L_{i,2} = 1, w_{i,2} = 1$
- The **adversarial approach** consists in building an **adversarial classifier** \hat{c} with target L and covariates X based on the concatenation of datasets $(X_{i,1}, L_{i,1}, w_{i,1})_{i=1,\dots,n_1}$ and $(X_{i,2}, L_{i,2}, w_{i,2})_{i=1,\dots,n_2}$
- Drift values** can be calculated by considering the **feature importance of the adversarial classifier**

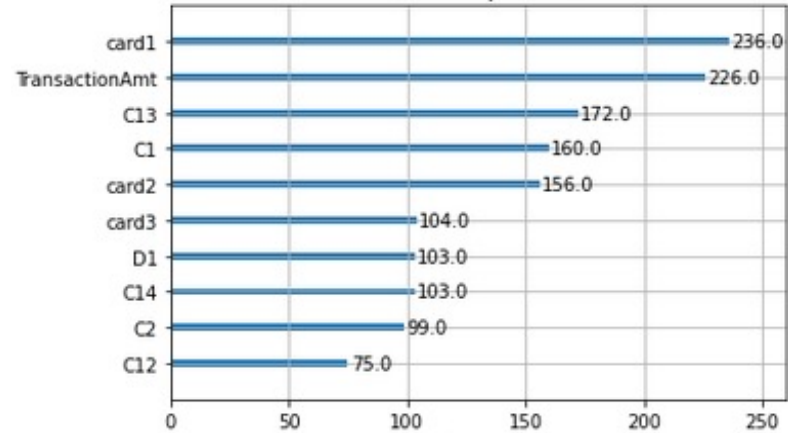


Adversarial approach: Example

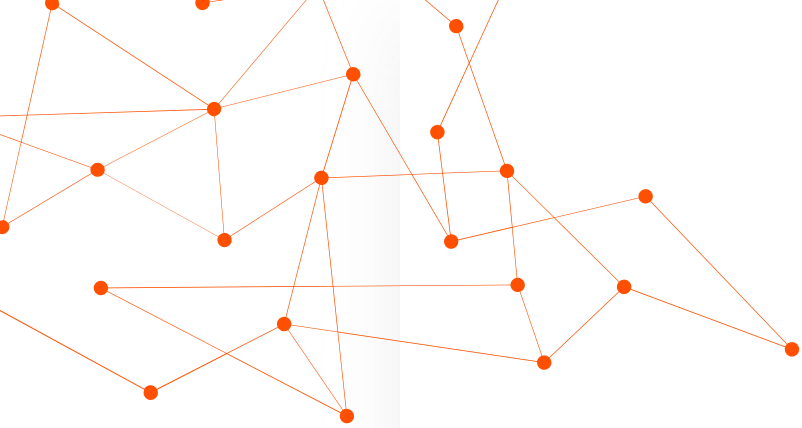
Drift values with the adversarial method



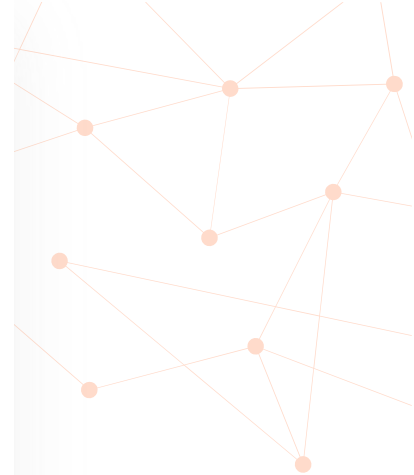
Feature importance of the model in production



- The drift values obtained with the adversarial approach may be combined with the feature importance of the model. **But it is hard to interpret !**



Correct data drift



Correct covariate shift with importance weighting

Property: Assume there is covariate shift between (X_1, Y_1) and (X_2, Y_2) and that $W(x) = \frac{dP_{X_2}(x)}{dP_{X_1}(x)}$ is defined.

Then for any function g :

$$E[W(X_1)g(X_1, Y_1)] = E[g(X_2, Y_2)]$$

proof: $E[W(X_1)g(X_1, Y_1)] = \iint W(x)g(x, y)dP_{Y_1|X_1}(y|x) dP_{X_1}(x) = \iint g(x, y)dP_{Y_2|X_2}(y|x) dP_{X_2}(x) = E[g(X_2, Y_2)]$

Morally, thanks to importance weights $W(x)$, it is possible to estimate the distribution (X_2, Y_2) based on the distribution (X_1, Y_1) .

- Idea

- Given the dataset $(X_{i,1}, Y_{i,1})_{i=1,\dots,n_1}$ and $(X_{i,2}, Y_{i,2})_{i=1,\dots,n_2}$, compute sample weights $W_{i,1}$ so that the distribution of $(X_{1,i}, Y_{1,i}, W_{i,1})_{i=1,\dots,n_1}$ approximates the distribution of $(X_{i,2}, Y_{i,2})_{i=1,\dots,n_2}$.
- Then train, validate, and select model with the dataset $(X_{1,i}, Y_{1,i}, W_{i,1})_{i=1,\dots,n_1}$ should result in better performance on $(X_{i,2}, Y_{i,2})_{i=1,\dots,n_2}$

Proposed procedure to correct data drift

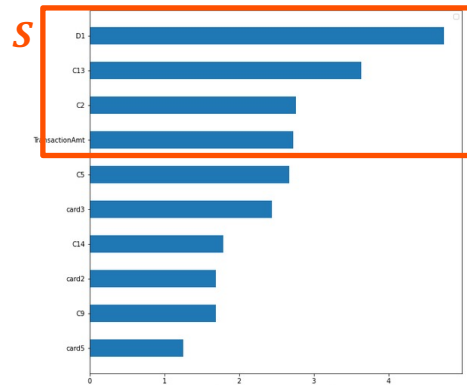
1) Based on model based drift values, **select a subset of input variables** $S \subset \{1, \dots, d\}$ for which we will correct the drift.

2) Use the adversarial approach to **build an adversarial classifier** \hat{c} **only based on feature** S . Let X^S the restriction of X to the variables in S .

3) **Compute weights** $W_{i,1} = \frac{\hat{c}(X_{i,1}^S)}{1 - \hat{c}(X_{i,1}^S)}$:

- Note that $W(x) = \frac{dP_{X_2}(x)}{dP_{X_1}(x)} = \frac{P(L=1|X=x)}{P(L=0|X=x)} = \frac{P(L=1|X=x)}{1 - P(L=1|X=x)}$ and thus $W(x)$ can be estimated by $\frac{\hat{c}(x)}{1 - \hat{c}(x)}$
- Cross validation is used in order to compute weights $W_{i,1}$

4) The dataset $(X_{1,i}, Y_{1,i}, W_{i,1})_{i=1, \dots, n_1}$ can then be use for model training, model selection.



Select subset of features S based on drift values



Demo



CinnaMon, a Python library for the monitoring of machine learning systems that focus on data drift: <https://github.com/zelros/cinnamon>





Conclusion



Conclusion

- Main ideas
 - Monitoring of the data drift
 - Focus on the data drift that has an impact on the ML system ?
 - Detection: Monitor indicators that makes sense from this perspective
 - Explain: Introduce drift values with a Model based approach
 - Correction of the data drift
 - Proposed methodology to correct covariate shift (using weights computed with adversarial approach) -> need more investigation
- Future work:
 - Closer next steps
 - Write an article about CinnaMon
 - CinnaMon library (add test, add documentation, etc.)
 - Research directions
 - Extend the model approach to other specific models
 - Model agnostic method to compute drift values (only relying on "model.predict" call)
 - Deal with streaming data (add .update() on top of .fit() method). In fact the question is how do you define D1 and D2)
 - Design a live alerting system (with robust alerts) based on the 3 data drift indicators
 - Benchmark the different ways to compute tree-based drift values

Thank you for your attention !



Bibliography

- 1) Sugiyama, Masashi, and Motoaki Kawanabe. *Machine learning in non-stationary environments: Introduction to covariate shift adaptation*. MIT press, 2012.
- 2) Sethi, Teggyot Singh, and Mehmed Kantardzic. "On the reliable detection of concept drift from streaming unlabeled data." *Expert Systems with Applications* 82 (2017): 77-99.



Zelros / [Z EH L R AO S]

Word invented by a recurrent
neural network trained on 130k
tech company names from
crunchbase.com



198 Avenue de France
75013 Paris
France



Balanstraße 73/Haus 10
81541 Munich
Germany



Corso di Porta Romana, 61
20122 – Milan
Italy

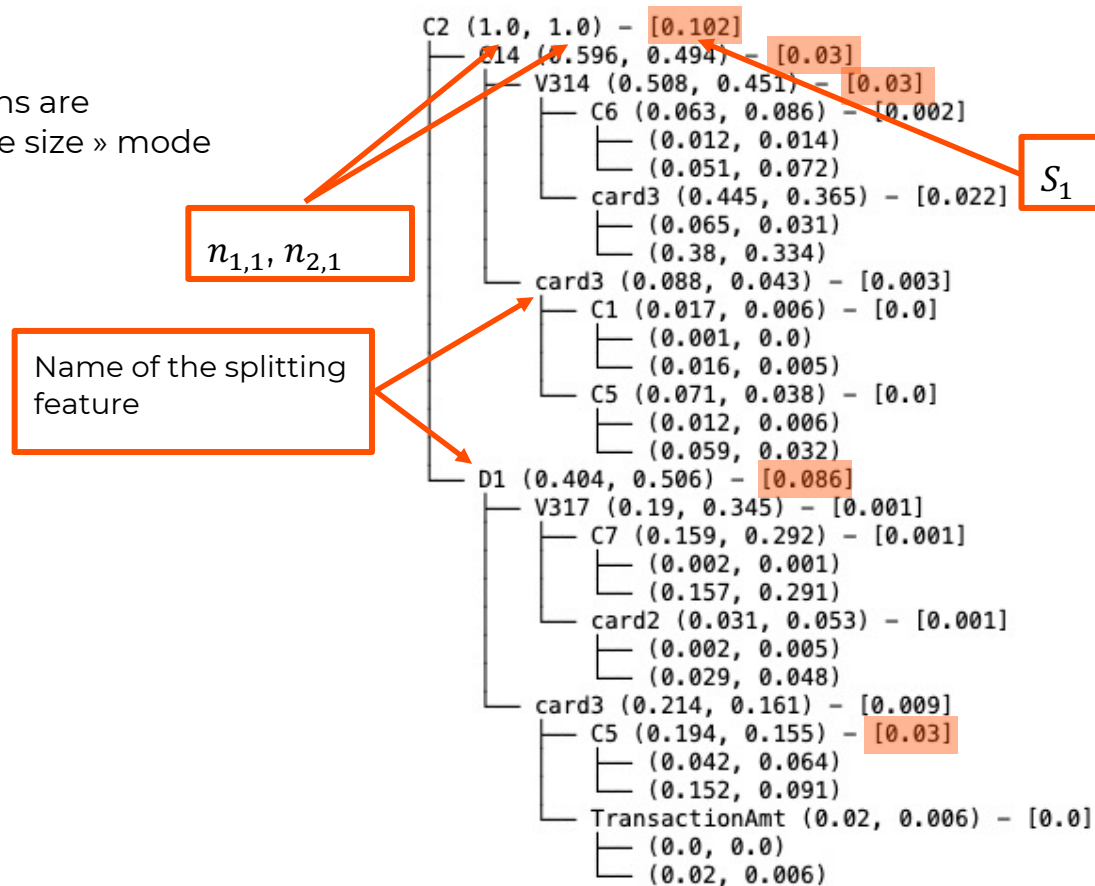


4 Place Ville-Marie, 2e et 3e étages,
Montréal, QC H3B 2E7
Canada



Drift values: illustration with tree

- Here split contributions are computed with « node size » mode



Adversarial approach (maths)

- Let (X, L) the random vector given by:
 - $L \sim \text{Bernoulli}(1/2)$
 - X follows the mixture distribution with:
 - $X|L = 0 \sim X_1$
 - $X|L = 1 \sim X_2$
- Let $(X_i, L_i, w_i)_{i=1, \dots, n_1+n_2}$ the concatenation of the datasets:
 - $(X_{i,1}, L_{i,1} = 0, w_{i,1})_{i=1, \dots, n_1}$ with $w_{i,1} = n_2/n_1$
 - $(X_{i,2}, L_{i,2} = 1, w_{i,2})_{i=1, \dots, n_2}$ with $w_{i,2} = 1$

Then $(X_i, L_i, w_i)_{i=1, \dots, n_1+n_2}$ follows the same distribution as (X, L)

- The **adversarial approach** consists in building a **discrimination model** \hat{c} with target L and covariates X (based on data $(X_i, L_i, w_i)_{i=1, \dots, n_1+n_2}$)
- Feature (drift) contributions can be calculated by considering the **feature importance of the discrimination model**